

# Evaluating Optimization Methods for Predicting Football Outcomes with Feedforward Neural Networks

Julius Graf                      Louis Sallé-Tourne  
julius.graf@berkeley.edu      louis\_salle-tourne@berkeley.edu

December 15, 2024

## Abstract

This academic project, inspired by the "Challenge Data - Football: Can you guess the winner? by QRT" competition, aims to explore the behavior of neural network learning processes, with a particular focus on the optimization algorithms used during training. The project is part of the UC Berkeley Fall 2024 graduate course INDENG 240: Optimization Analytics, taught by Prof. Phillip Kerger. By analyzing a dataset containing various information about football teams and their match outcomes, we will train neural networks to predict the winners of football matches. This will allow us to apply and evaluate the gradient descent techniques learned in class.

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Mathematical formalization</b>	<b>2</b>
2.1	Optimization problem . . . . .	2
2.2	Algorithms . . . . .	3
2.2.1	Gradient Descent . . . . .	3
2.2.2	Stochastic Gradient Descent . . . . .	3
2.2.3	Adam Optimization . . . . .	4
<b>3</b>	<b>Results</b>	<b>5</b>
3.1	Parameters . . . . .	5
3.2	Results . . . . .	5
3.2.1	Performance Comparison . . . . .	5
3.2.2	Convergence Plots . . . . .	5
<b>4</b>	<b>Conclusion</b>	<b>6</b>

## 1 Introduction

We aim to contribute to the 2024 "Challenge Data - Football: Can you guess the winner? by QRT"Football: Can you guess the winner? by QRT. Given that this project is part of the INDENG 240 course, our primary focus will be on the optimization aspects: defining the optimization problem, modeling through neural networks, and selecting appropriate optimization algorithms.

The competition organizers provide the training dataset for this challenge directly. Two tables give the input data:

- `train_home_team_statistics_df`: This table contains all information about the home team, with `game_id` as the key. It includes a variety of statistical data such as shots taken, passes made, and other performance metrics,

- **train\_away\_team\_statistics\_df**: This table has the same structure as the home team table but contains data for the away team, also keyed by `game_id`.

These tables can be joined using the `game_id` to create a comprehensive dataset for each match. The entire dataset comprises 12,203 rows, with 142 columns for each team (home and away). Given that we split the training data into training, testing, and validation data with an 80% ratio, the training set denoted  $\mathcal{D}_m$  is of size  $m = 7873$ . The columns are predominantly numerical, except for the first two: `LEAGUE` and `TEAM_NAME`. Thus, the input size is  $p = 280$  (dropping the league and the team name, one has 140 regressors for each of both teams). Examples of numerical columns include:

- **TEAM\_SHOTS\_TOTAL\_season\_sum**: Total shots taken by the team throughout the whole season,
- **TEAM\_SHOTS\_INSIDEBOX\_season\_sum**: Total shots taken inside the box by the team throughout the whole season,
- **TEAM\_SHOTS\_OFF\_TARGET\_season\_sum**: Total off-target shots taken by the team throughout the whole season,
- **TEAM\_SHOTS\_ON\_TARGET\_season\_sum**: Total on-target shots taken by the team throughout the whole season,
- **TEAM\_SHOTS\_OUTSIDEBOX\_season\_sum**: Total shots taken outside the box by the team throughout the whole season,
- **TEAM\_PASSES\_season\_sum**: Total number of passes made by the team throughout the whole season,
- **TEAM\_SUCCESSFUL\_PASSES\_season\_sum**: Total number of successful passes made by the team throughout the whole season,

The label dataset, `Y_train`, contains the same number of rows (12,203) and includes the columns `HOME_WINS`, `DRAW`, and `AWAY_WINS`, with a value of 1 indicating the match result. Therefore, the whole dataset size is 12,203. Our task is a multiclass classification problem (thus, the number of classes is  $c = 3$ ), where we aim to predict the outcome of each football match.

To achieve this, we will utilize the `PyTorch` library to build and train our neural network. This project will allow us to apply and evaluate the gradient descent techniques and other optimization algorithms learned in class. By focusing on the optimization aspects, we hope to gain insights into how different optimization strategies affect the performance of neural networks in predicting football match outcomes.

## 2 Mathematical formalization

### 2.1 Optimization problem

We consider a neural network  $f_\theta: x \in \mathbb{R}^p \mapsto W_2\sigma(W_1x + b_1) + b_2 \in \mathbb{R}^c$ , where the learnable parameters are  $\theta = (W_1, W_2, b_1, b_2)$ . These parameters belong to the space

$$\Theta = \mathcal{M}_{h,p}(\mathbb{R}) \times \mathcal{M}_{c,h}(\mathbb{R}) \times \mathbb{R}^h \times \mathbb{R}^c.$$

Here,  $W_1 \in \mathcal{M}_{h,p}(\mathbb{R})$  and  $W_2 \in \mathcal{M}_{c,h}(\mathbb{R})$  are weight matrices, while  $b_1 \in \mathbb{R}^h$  and  $b_2 \in \mathbb{R}^c$  are bias vectors. The activation function is set to  $\sigma = \text{ReLU}$ . For practical purposes, the parameter space  $\Theta$  can be assimilated into  $\mathbb{R}^d$  with

$$d = h(1 + p + c) + c.$$

This transformation is achieved by flattening the matrices  $W_1$  and  $W_2$  using a column-wise vectorization function  $\phi_{p,q}: \mathcal{M}_{p,q}(\mathbb{R}) \rightarrow \mathbb{R}^{pq}$ , defined as

$$\phi_{p,q}(A) = \begin{bmatrix} A_1 \\ A_2 \\ \vdots \\ A_q \end{bmatrix},$$

where  $A_i \in \mathbb{R}^p$  is the  $i$ -th column of  $A \in \mathcal{M}_{p,q}(\mathbb{R})$ . Thus,  $W_1 \equiv \phi_{h,p}(W_1)$  and  $W_2 \equiv \phi_{c,h}(W_2)$ . We adopt the cross-entropy loss function, defined as follows:

$$\forall \mathcal{D} \subseteq \mathbb{R}^p \times \mathbb{R}^c, |\mathcal{D}| < \infty, \quad \forall \theta \in \mathbb{R}^d, \quad l(\theta, \mathcal{D}) = -\frac{1}{m} \sum_{(x,y) \in \mathcal{D}} \sum_{j=1}^c y_j \text{softmax}(f_\theta(x))_j.$$

Here,  $\mathcal{D}$  is a dataset of finite size, with  $m = |\mathcal{D}|$ . For a specific dataset

$$\mathcal{D}_m = \{(x_i, y_i) \in \mathbb{R}^p \times \mathbb{R}^c \mid 1 \leq i \leq m\},$$

we aim to solve the following optimization problem:

$$(\mathcal{P}): \quad \min_{\theta \in \Theta} l(\theta, \mathcal{D}_m) + \frac{\kappa}{2m} \|\theta\|_2^2.$$

Here,  $\kappa$  is a regularization weight that penalizes the magnitude of the parameters. The focus of our study is to compare various optimization methods for solving the problem  $(\mathcal{P})$ .

## 2.2 Algorithms

In this subsection, we outline the optimization methods used to solve the problem  $(\mathcal{P})$ . These include Gradient Descent (GD), Stochastic Gradient Descent (SGD), and the Adam optimizer. Each method is described below.

### 2.2.1 Gradient Descent

The classic Gradient Descent algorithm iteratively updates the parameters  $\theta$  by following the negative gradient of the loss function.

---

#### Algorithm 1 Gradient Descent (GD)

---

- 1: Initialize  $\theta_0$ ,  $n \leftarrow 0$ ,  $\eta > 0$ ,  $\kappa \geq 0$  and number of epochs  $T$
  - 2: **for**  $n = 1, 2, \dots, T$  **do**
  - 3:    $\theta_n = \theta_{n-1} \left(1 - \frac{\eta\kappa}{m}\right) - \frac{\eta}{m} \sum_{i=1}^m \nabla_{\theta} l(\theta_{n-1}, \{x_i, y_i\})$
  - 4: **end for**
  - 5: **return**  $\theta_n$
- 

However, since the gradient is computed in a single run over the whole dataset, convergence will be pretty slow and the main bottleneck is memory inefficiency. We therefore next present two more advanced algorithms, that are also used in practice, to compare their performance afterward.

### 2.2.2 Stochastic Gradient Descent

Stochastic gradient descent works by studying the Markov chain  $\{X_n; n \in \mathbb{N}\}$  defined through its random mapping representation  $X_{n+1} = f_{n+1}(X_n)$  where

$$f_{n+1}(\theta) = \theta \left(1 - \frac{\eta\kappa}{b}\right) - \frac{\eta}{b} \sum_{i \in B_{n+1}} \nabla_{\theta} l(\theta, \{x_i, y_i\}).$$

Here,  $b$  is the (fixed) size of the mini-batch  $B_{n+1}$ , a random subset of  $\llbracket 1, m \rrbracket$ , dependent on  $n+1$  as it is reshuffled at each iteration. The functions  $f_n$  for  $n$  in  $\mathbb{N}^*$  are iid copies of a locally Lipschitz random mapping  $f: \mathbb{R}^d \rightarrow \mathbb{R}^d$ .<sup>1</sup>

---

<sup>1</sup>Yanlin Qu, Jose Blanchet, and Peter Glynn. *Deep Learning for Computing Convergence Rates of Markov Chains*. 2024. arXiv: 2405.20435 [cs.LG]. URL: <https://arxiv.org/abs/2405.20435>.

**Algorithm 2** Stochastic Gradient Descent (SGD)

---

```

1: Initialize  $X_0$ ,  $b \in \mathbb{N}^*$ ,  $k \leftarrow 0$ ,  $\eta > 0$ ,  $\kappa \geq 0$  and number of epochs  $T$ 
2: for  $n = 1, 2, \dots, T$  do
3:   Sample random permutation  $\pi$  according to  $\pi \sim \mathcal{U}(\mathfrak{S}_m)$ 
4:   Divide  $\mathcal{D}_m$  into  $N = \lceil m/b \rceil$  mini-batches  $B_j^{(n)} = \pi(\llbracket (j-1)b + 1, \min(jb, m) \rrbracket)$  for  $j$  in  $\llbracket 1, N \rrbracket$ 
5:   for  $j = 1, \dots, N$  do
6:      $X_k \leftarrow X_k \left(1 - \frac{\eta\kappa}{m}\right) - \frac{\eta}{b} \sum_{i \in B_j^{(n)}} \nabla_{\theta} l(X_k, \{x_i, y_i\})$ 
7:      $k \leftarrow k + 1$ 
8:   end for
9: end for
10: return  $\theta_k$ 

```

---

Note that when one chooses  $b = m$ , the algorithm becomes deterministic and corresponds exactly to the gradient descent algorithm. In fact, one will only have one mini-batch  $\pi(\llbracket 1, m \rrbracket) = \llbracket 1, m \rrbracket$  for all  $\pi$  in  $\mathfrak{S}_m$ . Now, as we've seen in class, momentum is a point that enhances convergence and renders the algorithm less prone to local minima. The following algorithm, Adam, goes beyond the idea of first-order momentum.

**2.2.3 Adam Optimization**

The Adam optimizer extends SGD by incorporating momentum and adaptive learning rates. It tracks the first- and second-order moments of the gradient using moving averages. Bias-corrected estimates  $\hat{u}_n$  and  $\hat{v}_n$  are computed to ensure stability. The update rule is:

$$\theta_k \leftarrow \theta_k - \eta \hat{u}_k \odot (\hat{v}_k + \epsilon^2)^{-1/2},$$

where  $\eta > 0$  is the learning rate,  $\epsilon > 0$  is a small constant for numerical stability, and  $\kappa \geq 0$  is the regularization weight.

**Algorithm 3** Adam

---

```

1: Initialize  $X_0$ ,  $u_0 \leftarrow 0$ ,  $v_0 \leftarrow 0$ ,  $k \leftarrow 0$ ,  $\alpha > 0$ ,  $\beta > 0$ ,  $\epsilon > 0$ ,  $\eta > 0$ ,  $\kappa \geq 0$ ,  $b \in \mathbb{N}^*$ , and number of epochs  $T$ 
2: for  $n = 1, 2, \dots, T$  do
3:   Sample random permutation  $\pi$  according to  $\pi \sim \mathcal{U}(\mathfrak{S}_m)$ 
4:   Divide  $\mathcal{D}_m$  into  $N = \lceil m/b \rceil$  mini-batches  $B_j^{(n)} = \pi(\llbracket (j-1)b + 1, \min(jb, m) \rrbracket)$  for  $j$  in  $\llbracket 1, N \rrbracket$ 
5:   for  $j = 1, \dots, N$  do
6:      $g_k \leftarrow \frac{1}{b} \sum_{i \in B_j^{(n)}} \nabla_{\theta} l(X_k, \{x_i, y_i\}) + \frac{\kappa}{m} X_k$ 
7:      $u_k \leftarrow \alpha u_k + (1 - \alpha) g_k$ 
8:      $v_k \leftarrow \beta v_k + (1 - \beta) g_k \odot g_k$ 
9:      $\hat{u}_k \leftarrow u_k / (1 - \alpha^{k+1})$ 
10:     $\hat{v}_k \leftarrow v_k / (1 - \beta^{k+1})$ 
11:     $X_k \leftarrow X_k - \eta \hat{u}_k \odot (\hat{v}_k + \epsilon^2)^{-1/2}$ 
12:     $k \leftarrow k + 1$ 
13:   end for
14: end for
15: return  $\theta_k$ 

```

---

## 3 Results

### 3.1 Parameters

The practical values used in this study are summarized in Table 1. Parameters are categorized into Model Parameters, Training Settings, and Optimization Settings for clarity. Mathematical notations are provided where applicable, along with the corresponding values.

Category	Notation	Value
<b>Model Parameters</b>		
Input Dimension	$p$	280
Number of Classes	$c$	3
Hidden Layer Dimension	$h$	32
Parameter Dimension	$d$	9091
Activation Function	$\sigma$	ReLU
<b>Training Settings</b>		
Learning Rate	$\eta$	0.001
Batch Size	$b$	64
Number of Epochs	$T$	20
Loss Function	$l(\theta, \mathcal{D})$	Cross-Entropy Loss
Regularization weight	$\kappa$	0.7873
Early Stopping Patience	–	5 epochs
Size of Training Set	$m$	7873
Train-Test Splitting Factor	–	0.8
<b>Optimization Settings</b>		
Optimizer	–	Adam, SGD, GD (comparative study)
Learning Rate Scheduler	StepLR( $\tau, \gamma$ )	StepLR (step size: $\tau = 20$ , decay factor: $\gamma = 0.5$ )
Betas	$(\alpha, \beta)$	(0.9, 0.999)
Numerical Stabilizer	$\epsilon$	0.001

Table 1: Parameters, mathematical notation, and corresponding values used for training and evaluation. Parameters are grouped into model-specific settings, training configurations, and optimization techniques.

### 3.2 Results

#### 3.2.1 Performance Comparison

The performance of the models trained using the SGD and Adam optimizers is summarized in the table below:

Optimizer	Final Validation Loss	Final Validation Accuracy (%)
SGD	1.0303	48.96
Adam	1.0690	48.84

Table 2: Performance comparison of SGD and Adam optimizers.

#### 3.2.2 Convergence Plots

The figures below show the convergence of the training and validation losses and the validation accuracies for both optimizers.

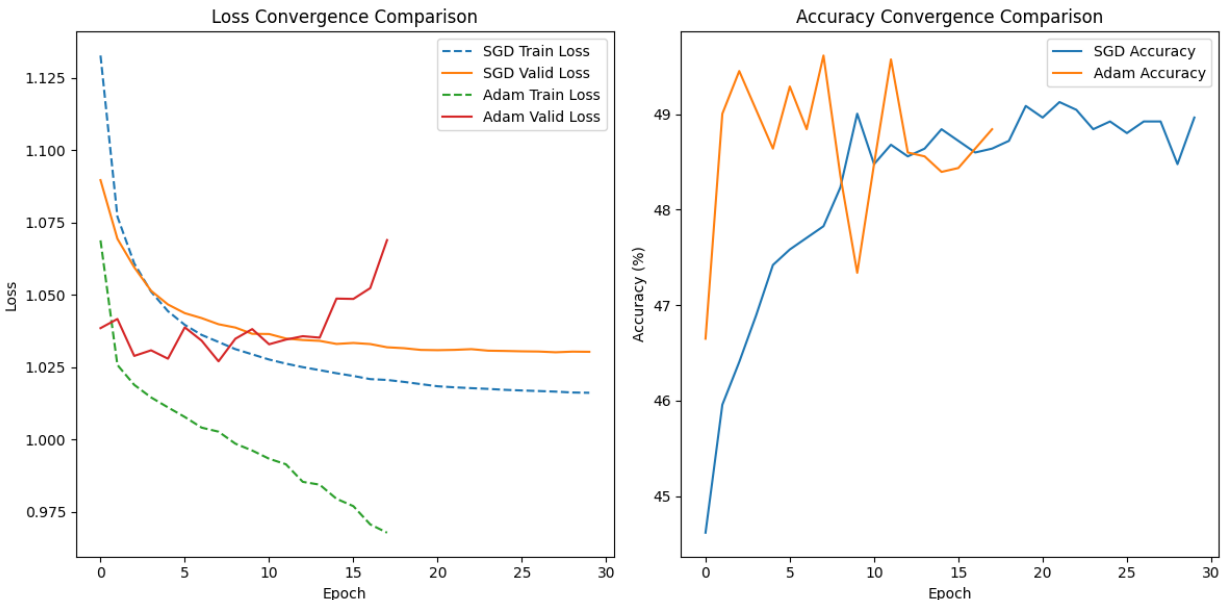


Figure 1: Left: Loss convergence for SGD and Adam optimizers. Right: Accuracy convergence for SGD and Adam optimizers.

## 4 Conclusion

This study evaluated the performance of feedforward neural networks trained with two optimization algorithms, SGD and Adam, to predict football match outcomes. Key results were compared against two benchmark models: (1) the baseline predicting that the HOME team wins, achieving 44% accuracy, and (2) a gradient boosting tree model trained solely on numerical features, yielding 47.5% accuracy on the training set.

The neural network models achieved validation accuracies of 48.96% with SGD and 48.84% with Adam, demonstrating marginal improvement over the benchmarks. Both optimizers displayed similar final performance but revealed notable differences in their behavior:

- SGD achieved slightly higher validation accuracy and lower final validation loss, indicating a more stable convergence to a solution that generalizes better. The simplicity of SGD's fixed learning rate may have contributed to more consistent updates during training, particularly for this problem's dataset and network structure.
- Adam, while competitive, showed a slightly higher final validation loss, which suggests a possible issue with overfitting or sensitivity to hyperparameter settings. Adam's adaptive learning rate mechanism can sometimes cause oscillations or premature convergence in scenarios where the learning rate becomes too aggressive or inconsistent. This could explain its slightly inferior performance compared to SGD in this study.

Despite these challenges, Adam exhibited faster initial convergence, as observed in the loss and accuracy plots. Its ability to adjust learning rates dynamically allowed it to make rapid progress early in training, which is a notable advantage over SGD. This aspect of Adam can be particularly beneficial in cases where computational efficiency or early stopping is critical.

In summary, while SGD proved marginally superior in this specific task due to its stability and generalization, Adam's faster convergence remains advantageous in certain contexts. The results underscore the importance of tailoring optimization strategies to the characteristics of the dataset and model. Future work could involve further hyperparameter tuning, regularization techniques, and hybrid training strategies to fully harness the strengths of both optimizers and achieve better overall performance.

## Appendix: Python3 Code

```

1  import pandas as pd
2  import numpy as np
3  from sklearn import model_selection
4  import warnings
5  import matplotlib.pyplot as plt
6  import torch
7  from torch.utils.data import DataLoader, TensorDataset
8  import torch.nn as nn
9  import torch.optim as optim
10
11 warnings.filterwarnings('ignore')
12
13 # Load and preprocess data
14 train_home_team_statistics_df = pd.read_csv('data/Train_Data/train_home_team_statistics_df.
15     csv', index_col=0)
16 train_away_team_statistics_df = pd.read_csv('data/Train_Data/train_away_team_statistics_df.
17     csv', index_col=0)
18 train_scores = pd.read_csv('data/Y_train_1rknArQ.csv', index_col=0)
19
20 train_home = train_home_team_statistics_df.iloc[:, 2:]
21 train_away = train_away_team_statistics_df.iloc[:, 2:]
22
23 train_home.columns = 'HOME_' + train_home.columns
24 train_away.columns = 'AWAY_' + train_away.columns
25
26 train_data = pd.concat([train_home, train_away], join='inner', axis=1)
27 train_scores = train_scores.loc[train_data.index]
28
29 train_data = train_data.replace({np.inf: np.nan, -np.inf: np.nan})
30 train_scores_indices = train_scores.idxmax(axis=1).map({'HOME_WINS': 0, 'DRAW': 1, '
31     AWAY_WINS': 2})
32
33 X_train, X_valid, y_train, y_valid = model_selection.train_test_split(
34     train_data, train_scores_indices, train_size=0.8, random_state=42
35 )
36
37 # Prepare data for PyTorch
38 X_train_tensor = torch.from_numpy(X_train.fillna(0).to_numpy()).float()
39 y_train_tensor = torch.from_numpy(y_train.to_numpy()).long()
40 X_valid_tensor = torch.from_numpy(X_valid.fillna(0).to_numpy()).float()
41 y_valid_tensor = torch.from_numpy(y_valid.to_numpy()).long()
42
43 train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
44 valid_dataset = TensorDataset(X_valid_tensor, y_valid_tensor)
45
46 train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
47 valid_loader = DataLoader(valid_dataset, batch_size=64, shuffle=False)
48
49 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
50
51 # Define the Neural Network
52 class NeuralNet(nn.Module):
53     def __init__(self, input_size, hidden_size, num_classes):
54         super(NeuralNet, self).__init__()
55         self.fc1 = nn.Linear(input_size, hidden_size)
56         self.relu = nn.ReLU()
57         self.fc2 = nn.Linear(hidden_size, num_classes)
58
59     def forward(self, x):
60         x = self.fc1(x)
61         x = self.relu(x)
62         x = self.fc2(x)
63         return x
64
65 # Define evaluation function
66 def evaluate(model, criterion, loader):

```

```

64     model.eval()
65     total_loss, correct, total = 0, 0, 0
66     with torch.no_grad():
67         for inputs, labels in loader:
68             inputs, labels = inputs.to(device), labels.to(device)
69             outputs = model(inputs)
70             total_loss += criterion(outputs, labels).item()
71             predicted = outputs.argmax(dim=1)
72             correct += (predicted == labels).sum().item()
73             total += labels.size(0)
74     accuracy = 100 * correct / total
75     return total_loss / len(loader), accuracy
76
77 # Early stopping implementation
78 class EarlyStopping:
79     def __init__(self, patience, verbose=False):
80         self.patience = patience
81         self.verbose = verbose
82         self.best_loss = float('inf')
83         self.counter = 0
84
85     def step(self, current_loss):
86         if current_loss < self.best_loss:
87             self.best_loss = current_loss
88             self.counter = 0
89             return True
90         else:
91             self.counter += 1
92             if self.counter >= self.patience:
93                 if self.verbose:
94                     print("Early stopping triggered.")
95                 return False
96             return True
97
98 # Train the model
99 def train_model(optimizer_name, lr=0.001, weight_decay=1e-5, hidden_size=32, num_epochs=30,
100                 early_stopping_patience=10):
101     input_size = X_train_tensor.shape[1]
102     num_classes = len(np.unique(y_train))
103
104     model = NeuralNet(input_size, hidden_size, num_classes).to(device)
105     criterion = nn.CrossEntropyLoss()
106
107     if optimizer_name == 'SGD':
108         optimizer = optim.SGD(model.parameters(), lr=lr, weight_decay=weight_decay)
109     elif optimizer_name == 'Adam':
110         optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=weight_decay)
111     else:
112         raise ValueError("Unknown optimizer")
113
114     scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=20, gamma=0.5)
115     early_stopping = EarlyStopping(patience=early_stopping_patience, verbose=True)
116
117     train_losses, valid_losses, accuracies = [], [], []
118     learning_rates = []
119
120     for epoch in range(num_epochs):
121         model.train()
122         train_loss = 0
123         for inputs, labels in train_loader:
124             inputs, labels = inputs.to(device), labels.to(device)
125
126             outputs = model(inputs)
127             loss = criterion(outputs, labels)
128
129             optimizer.zero_grad()
130             loss.backward()
131             optimizer.step()

```



```

131         train_loss += loss.item()
132
133     train_loss /= len(train_loader)
134     train_losses.append(train_loss)
135
136     valid_loss, valid_accuracy = evaluate(model, criterion, valid_loader)
137     valid_losses.append(valid_loss)
138     accuracies.append(valid_accuracy)
139
140     learning_rates.append(optimizer.param_groups[0]['lr'])
141
142     print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {valid_loss:.4f}, Accuracy: {
143           valid_accuracy:.2f}%, Learning Rate: {learning_rates[-1]:.6f}')
144
145     if not early_stopping.step(valid_loss):
146         break
147
148     scheduler.step()
149
150     return train_losses, valid_losses, accuracies, valid_loss, valid_accuracy,
151           learning_rates
152
153 # Plot metrics
154 def plot_metrics(train_losses, valid_losses, accuracies, title_suffix=""):
155     plt.figure(figsize=(12, 6))
156     plt.subplot(1, 2, 1)
157     plt.plot(train_losses, label='Train Loss')
158     plt.plot(valid_losses, label='Valid Loss')
159     plt.xlabel('Epoch')
160     plt.ylabel('Loss')
161     plt.title(f'Loss Convergence {title_suffix}')
162     plt.legend()
163
164     plt.subplot(1, 2, 2)
165     plt.plot(accuracies, label='Accuracy')
166     plt.xlabel('Epoch')
167     plt.ylabel('Accuracy (%)')
168     plt.title(f'Accuracy Convergence {title_suffix}')
169     plt.legend()
170     plt.tight_layout()
171     plt.show()
172
173 # Plot learning rate progression
174 def plot_learning_rate(learning_rates, title_suffix=""):
175     plt.figure(figsize=(6, 4))
176     plt.plot(learning_rates, label='Learning Rate')
177     plt.xlabel('Epoch')
178     plt.ylabel('Learning Rate')
179     plt.title(f'Learning Rate Progression {title_suffix}')
180     plt.legend()
181     plt.tight_layout()
182     plt.show()
183
184 # Plot comparison
185 def plot_comparison(sgd_results, adam_results):
186     plt.figure(figsize=(12, 6))
187
188     plt.subplot(1, 2, 1)
189     plt.plot(sgd_results[0], label='SGD Train Loss', linestyle='--')
190     plt.plot(sgd_results[1], label='SGD Valid Loss')
191     plt.plot(adam_results[0], label='Adam Train Loss', linestyle='--')
192     plt.plot(adam_results[1], label='Adam Valid Loss')
193     plt.xlabel('Epoch')
194     plt.ylabel('Loss')
195     plt.title('Loss Convergence Comparison')
196     plt.legend()

```

```
197     plt.subplot(1, 2, 2)
198     plt.plot(sgd_results[2], label='SGD Accuracy')
199     plt.plot(adam_results[2], label='Adam Accuracy')
200     plt.xlabel('Epoch')
201     plt.ylabel('Accuracy (%)')
202     plt.title('Accuracy Convergence Comparison')
203     plt.legend()
204
205     plt.tight_layout()
206     plt.show()
207
208 # Train and evaluate models
209 sgd_results = train_model('SGD')
210 adam_results = train_model('Adam')
211
212 # Plot results
213 plot_comparison(sgd_results, adam_results)
214 plot_learning_rate(sgd_results[5], title_suffix="(SGD)")
215 plot_learning_rate(adam_results[5], title_suffix="(Adam)")
216
217 # Final performance
218 print(f"SGD Final Validation Loss: {sgd_results[3]:.4f}, Final Validation Accuracy: {
    sgd_results[4]:.2f}%")
219 print(f"Adam Final Validation Loss: {adam_results[3]:.4f}, Final Validation Accuracy: {
    adam_results[4]:.2f}%")
```

---

Listing 1: See the repository [https://github.com/juliusgraf/challenge\\_qrt\\_optim](https://github.com/juliusgraf/challenge_qrt_optim).